The Mike Speiser Incubation Playbook - kwokchain

Kevin

36-45 minutes

In Formula 1 racing, you can win a world championship as a driver with one team but then not even make the top 10 without that team's car and infrastructure. Venture can often feel like this, too. Many top performing VCs would struggle if they weren't on their firm's platform. And similarly, a far greater number of VCs might be able to do well if they were just at a firm with a strong enough brand. Most special are those that are the source of their own success.

In Making Uncommon Knowledge Common, I wrote about Rich Barton because he's one of the rare founders (or investors) with the demonstrated ability to create multiple billion dollar companies. Unpacking and learning from the few who have shown repeatable and internally compounding approaches to building companies is important.

Unlike consumer, traditional enterprise markets lend themselves more naturally to deterministic and repeatable success. There's a small handful of VCs who have clearly shown they can succeed repeatedly and whose approaches and playbooks are legible enough to imply it's not a fluke. Speiser is one of them.

Speiser's portfolio includes companies like Pure Storage and Snowflake Computing. It's worth noting that Snowflake not only IPO'd and is now at a market cap of over $60B but Speiser and Sutter Hill Ventures owned more than 20% of the company leading up to the IPO. When Pure Storage went public, Sutter Hill held more than 25%. Speiser may have the highest percentage of portfolio companies that have become multi-billion dollar companies—and that trend looks to continue with his newer companies.

But impressive returns are not solely what matters for the industry. It's tempting to evaluate firms by their returns, and from the LP perspective that may be the correct metric. But another, and more important way to judge VC firms is by the value they add above replacement to their portfolio companies. How much do they help their portfolio companies increase their likelihood and magnitude of success? Firms do this most notably by providing capital, but also by other methods like lending their brand or directly helping with operations.

For founders, this value added is what matters. The returns of a VC firm only matter to a startup insofar as they translate into improved brand, network, or access to capital for the startup. A firm's financial performance is a reasonable signal that they may add real value and be worth partnering with, especially since some aspects like brand strength for recruiting, future financing, and customer development are a function of perceived firm success. But to prospective portfolio companies, a fund's returns are important only as a means, not an end.

What makes Speiser intriguing is how distinct his approach is from other VCs. The tantalizing clues suggest that he has figured something out that nobody else has: the formula for creating successful companies from scratch.

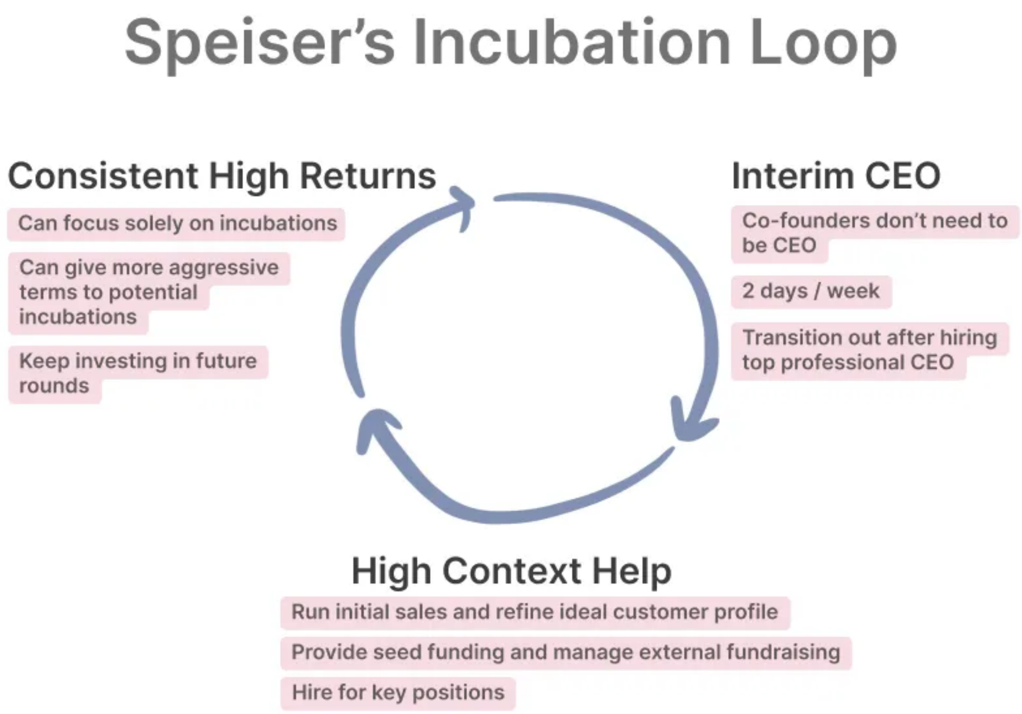

you didn't really think there wasn't going to be a drawing of a loop did you?

The Speiser playbook

At the core of Speiser's approach is incubating companies, or "originating companies" in Sutter Hill nomenclature. Instead of investing in existing companies, Speiser stays solely focused on one thing: starting and building companies. Even among others who have been very successful at incubations, he is the most singularly focused on this.

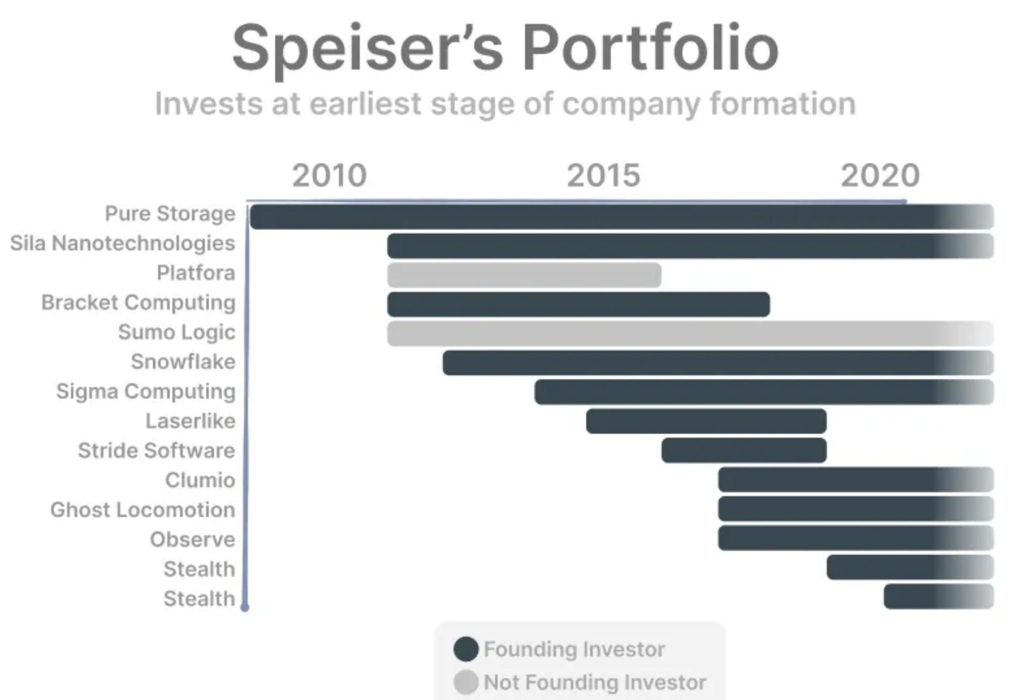

bespoke artisanal charts as a service

Every year Speiser incubates around one company. The core of his model is to find 2-3 co-founders and be the founding investor. Often he takes on the interim CEO role himself for the first year or two. This has many advantages. The biggest is that it reshapes the ideal founding team profile. He can focus on getting the right top technical co-founders that will have strong views on what to build and the ability to build it—even if they are people who don't generally view themselves as having a natural inclination to be founders. This is a significant talent arbitrage.

A better package for founders

There has been a dearth of coverage of Snowflake's three cofounders, Benoit Dageville, Thierry Cruanes, and Marcin Zukowski in both the news and social media. Partially this is because they have not sought the spotlight. But partially, it is due to the veneration of a certain type of founder we have, those who seek the limelight of public presence and being in control of every aspect of the company.

Snowflake's founders are cut from a different cloth. As Benoit Dageville put it "We never thought of it as building a company. We just wanted to build a cloud product. The company was an afterthought." Yet, their product and technical decisions have been prescient in threading the narrow path to taking on Amazon and Google in the most important core markets of cloud computing.

There are people who often don't *want* to be CEO, or even to start a company. They are driven by their conviction of what the future should look like, as well as their frustration with the internal dynamics they confront at legacy incumbents that prevent them from creating that reality. But they are still unlikely to start a company due to all the inertial cruft that comes with founding a company—and especially with being CEO. They want to build what matters, not set up a new corporate structure, manage fundraising, or build a sales team.

Eric Yuan, the founder and CEO of Zoom, has explained this feeling of being held back at Webex. He knew what should be built and that the Webex team could do it, but given the dynamics of Webex as a subsidiary of Cisco he was unable to get the political capital to do it. And he's proven himself right by leaving with his Webex teammates and building Zoom. Considering he was VP Engineering at Webex and still unable to build what he thought was important should be a very discomfiting reality shock to large companies about the very real economic harm the malaise of their internal processes have caused. However, for every Eric Yuan, there are countless others that never leave and start a company. The inertial barriers are too high. They can stay at their company and struggle to work on what they know should be built. Or leave and take on more uncertainty and risk than they want.

Speiser introduces a third model that breaks through this Scylla and Charybdis dilemma. Start a company with Speiser and stay focused on what you want: deciding what to build, hiring the team you need, and building it. Speiser will handle fundraising, handling the operations generally, and setting up the sales motion and machine. Founders get much of the autonomy and upside of starting a new company while also getting support and guardrails so they can stay focused while having confidence the business is being built well.

Speiser doesn't just take on these roles because founders don't want to do it. There are actually aspects of company building where he should be better than the founders. Sutter Hill Ventures has the capital already, so it's easy for them to take on responsibility for fundraising and remove that as a blocker. Instead of having to burn a lot of cycles fundraising, Sutter Hill can provide the capital. And they often do, leading multiple rounds into their companies. Or they can bring in outside investors, with the confidence that Sutter Hill can lead the entire round as a backstop if the process becomes too much of a hassle. Also, like any VC firm, Sutter Hill builds a brand that compounds their companies' ability to raise follow on funding. At this point there are multiple firms that have made their bread and butter following on after Sutter Hill, to great success.

Similarly, Speiser is likely to have more experience in setting up companies and the initial customer development process than the founders will. Perhaps most importantly, he has relationships with customers and an established reputation that can be used to bootstrap the initial pilot conversations, which may be the point of highest leverage for these new startups.

These advantages all *compound* with every incremental company Speiser originates, and not just because of the typical brand network effects that venture has broadly. In many tangible ways, the spread between Speiser's process knowledge relative to a new founder should widen with every new company.

As an industry we seem to often want to see machismo and martyrdom in founders. A decade ago it was wanting founders to be willing to mortgage their house and their kids' college fund. Now it is founders wanting to be in charge of every aspect of companies. If founders aren't willing to put everything on the line for the company their companies will be worse is the thought. As an ecosystem it doesn't appear the data bears this out. Everything we do that has expanded opportunity and decreased the friction to more people becoming founders has led to huge benefits for the industry.

Just as Eric Yuan should be a massive shot across the bow for all large tech incumbents, Snowflake's founders should be a wakeup call to venture that we have much further to go to enable and support even more brilliant people who don't think of themselves as CEOs to bring their vision of the world into existence.

A better package for CEOs

But this isn't the only talent arbitrage Speiser's playbook benefits from. The interim CEO model allows another one too. As the startup does well and figures out its product market fit, Speiser eventually rolls off as CEO and finds a full-time replacement to take on the role as he takes a step back into being solely a board member.

His companies are *very* advantaged in finding great CEOs to take the mantle. To understand why, look at it from the perspective of an executive looking to become the CEO of a company. Like the potential founders, these executives have their own Scylla and Charybdis dilemma. They want to be CEO of a company, but they also want to join a company that has already derisked product-market fit instead of founding a new startup with all the attendant risks. However, there is significant adverse selection among Series A and B companies that are looking for an external CEO.

Now that most founders want to stay on and scale up as CEO of their companies, it often indicates the company is struggling if the board is looking to replace the founder. Even if turning the company around is doable, the internal and board dynamics are likely to be very acrimonious—with a hostile deposed founder.

There are a few exceptions, like Linkedin and Hashicorp, where the companies were doing very well and the founders wanted to bring on a CEO. That can be extremely effective, with the founders and CEO partnering to great effect. There is a lot to emulate in the dynamic and shared responsibility between Reid Hoffman and Jeff Weiner at Linkedin or David McJannet, Armon Dadgar, and Mitchell Hashimoto at Hashicorp.

But increasingly, founders wanting to bring in a CEO are the exception, not the norm, if companies only look for new CEOs while pushing out the founder.

Speiser can offer a much more attractive package.

When Speiser talks to a potential CEO he can say he's showing them his strongest companies, because that's true. His model is built on him leaving his role as interim CEO once they are working—and moving on to do it over again at a new incubation. Finding a new CEO is a feature of success, not a cry for help.

Many of my friends who are executives are bombarded by VCs trying to trick them into taking on their worst performing companies. After they do some light diligence and referencing, they realize that the company is months away from failure or the founder will be hoping for them to fail from day one. Executives grow to realize they should be default skeptical of any companies that VCs try to recruit them for as CEO.

Speiser is one of the few VCs who will really be pushing you towards his best companies. This gives him a huge advantage in building his relationships with the best executives, because they know he will actually be helping them find companies they'd *want* to be running.

And the results support this. Take Snowflake Computing as an example. After Speiser stepped down, Snowflake brought on Bob Muglia as CEO. Muglia, who was previously EVP of Software at Juniper Networks and before that President of Servers and Tools at Microsoft, was an astounding get as CEO for a two year old startup. And then last year, Snowflake brought in Frank Slootman, formerly CEO of ServiceNow and DataDomain, who may be the greatest enterprise CEO of his generation to run Snowflake.

Potential downsides of model

However, Speiser's model is not without tradeoffs.

For prospective founders, equity is one example of this. While VCs like incubations because they are able to command higher ownership percentages, this comes at the cost of founders ending up with less ownership. For founders who would start a company no matter what and don't feel like they need much support, an incubation model like Speiser's will leave them with significantly less equity than they could get otherwise. Benoit Dageville, the Snowflake co-founder with the most equity, had 3.4% at IPO—less than founders often have at IPO.

For many founders, this tradeoff will be worth it after weighing how much Speiser and Sutter Hill increase the probability and scale of potential success. After all, Benoit's stake is currently worth over two billion dollars. But for many, it may not be worthwhile.

Another tradeoff is autonomy. Just as founders and later CEOs may want Speiser's experience and hands-on help building the company, they may regret his involvement where they differ in viewpoint. And the tradeoff of a VC being intimately involved in the business…is that they are intimately involved in the business. This is especially true where company and individual incentives may differ. Bringing Frank Slootman in as CEO, for example, could only come with the departure of Bob Muglia. I'm sure Muglia would have preferred to stay CEO, but the board decided they couldn't pass up the opportunity to bring in Slootman. With two billion dollars in equity, Muglia may do it all again even knowing that. But there is a tradeoff in autonomy that comes with incubating a company with an investor so closely involved.

There are real tradeoffs for investors in the model as well. Incubating a company, especially as interim CEO, is significantly more work than only investing in a company. Which means an investor like Speiser can make very few investments, so the cost if any of them don't work out is higher. By only doing incubations, an investor also misses out on being able to invest in promising companies and teams that are already formed. Finally, by primarily incubating companies, it can provide some misalignment with existing founders that may be more reticent to meet since the investor is unlikely to invest but may incubate a competitor.

Supported by structure and process

Speiser's incubation model breaks many conventions of and assumptions about venture. They even defy the conventions of how most firms do incubations. Spending two days a week with companies consumes more time and focus than most VCs do. Speiser does fewer investments, which only works because of an implicit assumption that his companies have a much higher likelihood of success. Speiser currently has a roughly 20% hit rate of his companies achieving multi-billion dollar valuations. Comparatively, top VC firms are typically at single digit percentages. The ratio of other top firms is lower because they make more investments, but nevertheless Speiser's hit rate is exceptional. His structure and process are integral to this. And similarly these success rates are what allow him to double down on his model. But how does he choose which spaces and companies to focus on with his scarce slots?

Underlying Speiser's approach is a belief that ideas matter. And you can make success way more than most believe.

Technical risks and secular shifts

The recurring core of Speiser's approach is finding a market undergoing a massive secular shift—and betting all-in on the full transition. He favors companies where the demand if it works is high—but the technical risk is high and most don't believe it can be done yet. All investors like benefitting from market tailwinds, but it is a very different thing to bet on one well before it has manifested and to the exclusion of more conventional and existing approaches.

When Snowflake Computing was started in 2012, most investors and companies were convinced that in order to sell into large companies you had to support on-premise workloads. Hell, most *customers* were convinced to sell to them you had to handle on-premise. Betting all-in reminds me of something Reid Hoffman once mentioned about principles. Principles are only principles if you'd hold them even when they are costly. It is only betting fully when it comes at a real cost to the business. For example, in customers that can't be served because they need a hybrid solution.

Ghost Locomotion, another of Speiser's incubations, is another example of this. It's a purely computer vision based approach to autonomous vehicles (AV).

In AV most companies take a hybrid computer vision and LiDAR approach. Google's Self-Driving Car team and the many teams started by Google SDC alums (Aurora, Nuro, Argo, Otto) use this hybrid approach. Though LiDAR is expensive, the thesis is that for a use case as high stakes as autonomous driving it's important to have heterogeneous sensors that can make up for each other's shortcomings. Historically these sensor fusion approaches have outperformed those focused solely on computer vision. Of course, general consensus is that on a long enough time scale computer vision alone will work for autonomous driving. After all you and I both use a computer vision approach and are able to drive. The key question is which approach will be first to market with a system that is accurate enough.

The bet for most autonomous vehicle companies is that the progressing along the current hybrid approach will work, and that teams that have worked in autonomy will have an advantage—having seen the ceilings on performance that need to be overcome.

Ghost Locomotion, like Tesla, has a very different approach (it should be noted that Tesla also is a coupled bet on the feedback loop of large scale car telemetry data being key). They only use computer vision and machine learning—a purely software approach, rather than a hybrid hardware and software one. It is a bet that machine learning will improve at a fast enough rate to surpass hybrid approaches, or perhaps that the status quo approach cannot reach sufficient accuracy at all. It's an aggressive view that prior domain expertise in autonomous vehicles is less important than expertise in AI and software engineering.

What approach will win in autonomous vehicles is entirely out of scope of this essay, but I point this out to say that going all-in on secular transitions is hard. It involves very real tradeoffs that will feel like they are wrong for a while and *are* often just wrong.

Going all-in on a market transition actually requires a *more* precise viewpoint on timing. Companies that do so rarely die because the market transition never happens—they die because the secular shift doesn't happen fast enough. Understanding when the market is ready and how to help catalyze the transition is key.

Speiser explicitly seeks these secular shifts in the companies he incubates. Taking on technical risk over distribution risk. It's the most publicly prominent feature of his investments. And you'll see it throughout his thoughts wherever they are documented.

Snowflake Computing is all-in on cloud data warehouses. Not hybrid or on-premise data warehouses.

Pure Storage is all-in on flash storage. Not hybrid or disk based storage.

Ghost Locomotion is all-in on computer vision driven autonomous driving. Not hybrid or primarily LiDAR based autonomous driving.

Speiser's companies take an aggressive view on transitions in the market, seeking out shifts that would create obvious and differentiated value that incumbents can't provide, but that many don't think will come yet. And Speiser's companies go all in.

They are betting that the shifts will come sooner rather than later. Or perhaps more importantly, that they can *be* the catalyst that accelerates the market flipping abruptly towards the future.

Speiser is willing to underwrite the technical risk for years. Like many firms, Sutter Hill Ventures has enough capital to support a new company for years. But where most firms would balk, Speiser will continue to invest on that bet that the market will catch up to his view. And he has demonstrated repeatedly a willingness to continue to plow millions more into follow-on rounds of his incubations, even while they are pre-traction. Sutter Hill led the first two rounds in Snowflake, and continued to invest more money into later rounds.

Built-in rigor

As interim CEO, Speiser can bring his own rigor to the entire process from idea conception through finding product-market fit and being set up for scaling. This rigor can be seen even before a company or approach is solidified. My understanding is that Speiser meets with hundreds of potential founders and customers before deciding who to incubate a company with. This is orders of magnitude more than most VCs who meet with a handful of candidates when deciding on an incubation.

The rigor continues once the company is formed. Speiser leads the customer development process as CEO. This is key because in the early days of a startup, sales isn't about revenue. It's about product and market discovery, so a tight feedback loop is needed between sales and product. How to run this initial enterprise sales motion is very different from normal sales—and something few have experience in. But Speiser has lots of experience since this is the stage he repeatedly focuses on. These initial customer pilots are also much easier with a strong brand and pre-existing connections to potential customers. Through prior investments, Speiser has more relationships with potential customers to call upon and a more refined playbook for the steps of honing in on the ideal customer profile, how to structure the pilots, and more.

Most VCs fall into an uncanny valley. They won't do work directly, so founders can't offload work (and the cognitive load that goes with it) off their plate to their investors. But they don't have enough context on a frequent enough basis to be able to really help shape the meta-process towards success. Speiser has both as interim CEO.

Considerations for the venture industry

Focus

Fittingly, Speiser's approach to venture is the same as his approach to other markets: never go hybrid.

We see this in his decisions about structuring his work. Besides his investment in SumoLogic's Series B, it appears he hasn't made any investments beyond his incubations. There are other VCs in Speiser's cohort with similarly impressive track records with incubations, such as Jim Goetz at Sequoia or Asheem Chandra at Greylock. But Speiser is rare in now *only* doing incubations.

At a first glance this appears odd—after all, there are many great companies that he didn't incubate that his reputation could help him get access to investing in. Isn't he sub-optimizing returns by not doing traditional investing?

The question is whether the loops of incubation and venture investing are additive to each other. In broad strokes they clearly are. After all, many of the same people who could be good founders to incubate companies with would be good founders to invest in separately, too. Much of the understanding of markets, company progress, and more are generalizable between the two approaches as well.

However, on deeper inspection they are actually not that aligned in many ways. Companies are a sequence of de-risking functional loops. Incubations are focused on the very earliest stages of these. By the time a traditional VC firm invests, companies are already set on many of the aspects that someone incubating companies must be good at.

While there is overlap in networks, in many ways there is also a tradeoff between what's best for incubations vs. investing. By focusing solely on incubations, Speiser does not need to keep pace with the torrent of potential investments that occupies most investors' schedules. This frees up significant amounts of time and focus.

Like a gas, the normal flow of investing can expand to fill an entire schedule. There is always a company of the week that must urgently be pursued. It's very hard to carve out real time for incubations amid this.

While other investors incubate companies, there is little differentiation between how they handle their incubations and their traditional venture investments. Each part of Speiser's process is tuned specifically towards incubations. From deciding which parts of the company he'll run for first few years to the talent arbitrage. Or his process for finding founders and spaces to incubate companies around. And from changes he's made to the incubation process over the last few years, it's clear that he continues to work refining the process. Going forward it will be interesting to see 1) whether he can scale the number of companies he can incubate and remove himself as the bottleneck to incubations 2) if he can create more sources of value to incubated companies that compound further with every incubation he does.

Optimal firm structure is downstream of the expected value distribution of portfolio companies. If incubations have a different likelihood and magnitude of success than traditional venture portfolios, it's inevitable the firm structures that best support them will also be different.

Visibility and brand strategy

Speiser is one of the top current VCs by returns, and after the Snowflake IPO he'll have likely returned more than $12B in returns to Sutter Hill. That will bring the average billion dollar plus outcome in his portfolio to around 20%, and likely to rise with more companies like Sigma Computing, Clumio, and Ghost Locomotion approaching that valuation.

More important, he is one of the few with a unique and clear strategy that has clear compounding loops. Why then is he virtually unknown in Silicon Valley at large? Especially when compared relative to other investors with similar track records.

Many firms and investors rely on broad top-of-funnel brand network effects. Andreessen Horowitz's model grows stronger as more people are familiar with Andreessen Horowitz. It's their brand network effect that drives dealflow, name recognition, and improves their portfolios' cost of recruiting or customer acquisition.

However, Speiser's model doesn't rely on these brand network effects optimized for inbound, so he doesn't have similar pressure to broadcast his strategy and success. In fact, he avoids it.

Firms in this quadrant are most interesting to study. Firms with strategies that require mainstream awareness are much better understood—because their approach is out in the open. But firms that are successful but don't rely on a broad brand network effect are much less understood—so there is greater potential they have discovered an un-arbitraged source of compounding returns.

> This is graphic I made for myself. You can tell because it's worse than my normal graphics. Bet you didn't know that was possible

That's not to say he doesn't have brand network effects at all. Within the network of people he'd want to start companies with, he is well known—but there's no benefit to him in being known more widely. The percentage of his meetings that are outbound is likely significantly higher than most VCs, so what matters is that if he reaches out to people they view him as credible.

Speiser's model is far more interesting than most firms', because it bucks the current strain of conventional venture thought. It throws our assumptions of how companies must be started into disarray by abolishing the strict Chinese firewall typically held between VCs and the operations of their portfolio companies. It believes venture firms can increase the likelihood of company success much more than people think, which runs counter to the current view of picking and access as the primary frontier of competition.

Most importantly, it understands that the function of venture is more important than its form. And it adapts its structure to best serve its purpose—improving the fundamental derisking of companies.

Future of incubation

For the most part, the venture industry is skeptical of its ability to incubate successes with any consistency. This pessimism can be seen implicitly in discussions of venture success, which revolve primarily around the knowing of and getting access to the top fundraises. And can be seen in the structure of most firms' portfolios and follow on investments

Speiser's success incubating, along with a few others, is important because it is an existence proof of venture's ability to meaningfully impact companies. Over time industries tend to get overly narrow, focusing on approaches that are known to work well. This is natural, but often means that without outliers that take new approaches and perturb the equilibrium, industries can get stuck on local maxima—with no one wanting to be the first to try unproven strategies.

Speiser's success ratio is far beyond the normal distribution of venture outcomes, and the fact that they are incubated with him as interim CEO indisputably proves his involvement is effective.

Will the magnitude of success of Snowflake's IPO trigger a re-examination of incubations? It should. With caveats.

Speiser's model is in stark contrast to modern norms around the optimal level of investor engagement with companies. The Chinese firewall around investor and company engagement is primarily a function of fear of investors harming the companies—whether maliciously or more often unintentionally. However, this should mean that there are high returns where there is trust between founders and investors on greater levels of context and engagement by investors.

And as a market, it's important to have high returns on increased trust, otherwise trust becomes purely an aesthetic attribute.

More importantly, the power dynamic in tech is shifting very strongly away from investors to founders. As this trend continues and founders increasingly are less afraid of their investors, the field will increasingly bifurcate. Investors will either be less engaged and more passive (primarily contributing capital and brand) or they will engage closer with companies but be held to a higher bar of helping push progress for the companies. Over time as the leverage moves towards founders and the fear of investors being able to damage companies continues to fall, there will be more openness to new formats of investor and company engagements.

Incubations have far more degrees of freedom than investing, since investors are more closely involved in many more aspects of the company—especially at the proto-formation stage before many core decisions have calcified. If you look across the current landscape, there are very different strategies across the taxonomy of firms that incubate. Some of the investors that incubate include Asheem Chandra (Greylock), Aneel Bhusri (Greylock), Jim Goetz (Sequoia), Kevin Ryan, SciFi VC (Max Levchin's fund), Thrive Capital, BoldStart VC, 8VC, Unusual Ventures, and many others. Just looking at this list there are many axes they all differentiate on in approach, what they think should be centralized by the incubating firm, and where they think value is generated. And the types of companies they incubate and the dynamic range of their outcomes is equally wide.

People often bucket all incubations together as a category. This is wrong as there are more variants of incubations than of traditional venture. VCs and LPs trying to do incubations without a clear viewpoint on how to press the form and structure of their approaches to best incubate will be disappointed.

In the coming year we are likely to see many moves to incubate companies. This will not be due to a change in the efficacy of incubating, but rather the shifts in market valuations—and the sharp rise in demand and valuations for earlier stage companies in certain segments.

When companies at $1M ARR can regularly raise in the hundreds of millions of dollars, the pull to invest earlier grows. Competition to invest in the companies attracting these multiples has grown so fast that investors are moving aggressively earlier and earlier in their lifecycle. In some ways this is rational as the industry has improved at inferring forward revenue predictability of companies, but in large part investors are moving to invest before all the data points on go-to-market have become clear. But as seed investing eventually also becomes very competitive and investments in new startups get marked up at high valuations, investors will increasingly look towards incubating new startups.

At some stage the constraint becomes the number of companies being founded in these areas. Incubating new companies will be one of the few ways for investors to generate proprietary deal flow.

Final thoughts

As a community, it's easy for tech to become reflexively fixated on returns. Especially in this environment where every day brings a new company raising at astronomical valuations or going public at $75B, it's easy to get lost in the allure of it. But lost in all of the breathless discussion is what is actually being built, what value is being added, and whether our ecosystem is improving.

Returns untethered from value creation are a temporal anomaly. Over a long enough time period, returns should accrue to where value is created. And we should be most worried when returns don't have any tie to value created—an ecosystem with no fitness function cascades to the worst disasters.

Just as companies are judged on their net present value of future cash flows, so too should we judge organizations on the net present value of their future value added.

Have you ever seen a sadder portrait of the rise and fall of empires than this drawing?

It’s easy to judge VC firms on returns. But value created above replacement is the much more interesting and leading metric, especially from a founder’s perspective. If the Ghosts of Christmas removed a VC firm from existence, how would all of its now former portfolio companies fare? Would they have raised from another VC firm and been in the exact same place? Or would they not even have existed?

At an ecosystem level this same principle applies. Companies and firms push progress by taking novel approaches that others can all learn from. Increasingly firms have converged on structure and strategy—a local maxima. Few firms perturb the equilibria. Often they fail. But it’s the pursuit of new approaches to every aspect of company building that perturb the equilibria and push progress. It’s these aspects of companies that create positive externalities for everyone.

Speiser’s approach breaks from venture convention. And with Snowflake’s recent IPO, it’s getting attention that’s long due.

In many ways, Speiser’s playbook reminds me of TikTok. The most beautiful aspect of TikTok’s business model is how every aspect of it is aligned with its network effect. Its short-form video format decreases friction and maximizes volume to feed into its network effect. Its algorithms favoring of shares and full watch-throughs are core metrics for engagement and distribution. Its duet and reply videos turn every aspect towards the creation of new videos. Its sounds, memes, and dances are all user-generated social capital marketplaces. It’s glorious to see a company willing to rethink everything to best make its form fit its function.

Speiser’s incubation model shares that intentional design and craft. Over time, each aspect of his model has been shaped to better serve originating companies—rather than just repurposing the traditional venture model for a new function.

And it’s not lost on me that similar to the talent arbitrage I discussed Union Square Hospitality Group having in Aligning Business Models to Markets, Speiser also has one in attracting EIRs and CEOs.

Among my friends who admire Speiser’s model, one question hangs over us all. Is Speiser’s model repeatable, or is Speiser unique? Can his playbook be improved and scaled, or is it a craft that is fundamentally artisanal? Have we found a new frontier of Coasean logic to company formation that will become mainstream with time, or a single bloom that will vanish with him?

There’s reason to be optimistic. While Speiser feels like an outlier in venture, there are a number of individual partners that have incubated companies with models that feel very similar to great success. It is the intentional refinement of structure and process that seem unique—rather than the entire thesis.

Over the coming years we'll hopefully see many investors evolve their own perspectives on Speiser's approach and execute on it with their own modifications. There are many iterations that I'm particularly interested in seeing executed.

Of course, what excites me most about Speiser is that he is not done. In recent years, it's clear he's continued to push experiments on how to further refine the incubation model: how to scale it up while centralizing and compounding the value the firm provides and how to use the unique advantages of his model and financial returns to provide a more compelling package to founders while building in defensibility. And of course, also failed experiments in where the model can be expanded to.

I love what Speiser has done and continues to work toward. The true next dream in my opinion is scaling and systematizing his playbook beyond being constrained primarily by his individual effort. This will be hard—perhaps even impossible if the core of his success turns out to be his connections or some ineffable sense of taste.

But if it's doable, it would be one of the most important developments in Silicon Valley and tech. His model is a better abstraction layer and structure for creating successful companies and it's a process that compounds in how effective it is with every company. It's ultimately a better way to truly drive innovation faster, rather than being merely a tollbooth on innovation.

Acknowledgements

Many thanks to Keila Fong, for the many discussions about this topic and help with this piece.

Additionally, thanks to Max Bulger, Michael Dempsey, Casey Winters, Saam Motamedi, and Zach Brock, for their discussions, edits, and help with this piece.